

Class 7: Machine Learning 1

Lilia Jimenez (PID:A16262599)


2024-01-30

#CLUSTERING METHODS The broad goal here is to find groupings (clusters) in your input data

##Kmeans

first, lets make up some data to cluster.

```
x<-rnorm(1000)
hist(x)
```



Make a vector of length 60 with 30 points centered at -3 and 30 points centered at +3

```
tmp<-c(rnorm(30, mean=-3), rnorm(30, mean=3))
```


I will now make an x and y dataset with 2 groups of points

```
x<- cbind(x=tmp, y=rev(tmp))  
x
```

	x	y
[1,]	-0.8954690	2.2379458
[2,]	-1.9822267	4.9903837
[3,]	-2.2306257	2.5003111
[4,]	-3.5185138	4.2328734
[5,]	-2.1666392	3.3579368
[6,]	-4.9207311	4.3877801
[7,]	-3.0210995	3.5407611
[8,]	-1.1608206	3.8203228
[9,]	-1.7126827	4.3740536
[10,]	-3.4617375	2.7849770
[11,]	-1.7585610	2.8464933
[12,]	-2.1312152	4.0897340
[13,]	-2.8819893	1.3263860
[14,]	-2.9600204	3.5382516
[15,]	-3.9553219	3.3458462
[16,]	-2.0976765	4.5494848
[17,]	-5.3111815	5.1389394
[18,]	-3.5621290	2.1757020
[19,]	-1.2471528	1.9815847
[20,]	-3.4998367	0.4017925
[21,]	-2.2425046	3.6576809
[22,]	-2.4088333	3.5906341
[23,]	-3.1534786	4.0272653
[24,]	-1.9640569	4.3402315
[25,]	-2.8621716	2.2968796
[26,]	-3.3953630	2.0108646
[27,]	-1.6651982	2.9713095
[28,]	-3.8819205	3.2694195
[29,]	-1.8740991	3.0180333
[30,]	1.0093213	1.8378191
[31,]	1.8378191	1.0093213
[32,]	3.0180333	-1.8740991
[33,]	3.2694195	-3.8819205
[34,]	2.9713095	-1.6651982

```
[35,] 2.0108646 -3.3953630
[36,] 2.2968796 -2.8621716
[37,] 4.3402315 -1.9640569
[38,] 4.0272653 -3.1534786
[39,] 3.5906341 -2.4088333
[40,] 3.6576809 -2.2425046
[41,] 0.4017925 -3.4998367
[42,] 1.9815847 -1.2471528
[43,] 2.1757020 -3.5621290
[44,] 5.1389394 -5.3111815
[45,] 4.5494848 -2.0976765
[46,] 3.3458462 -3.9553219
[47,] 3.5382516 -2.9600204
[48,] 1.3263860 -2.8819893
[49,] 4.0897340 -2.1312152
[50,] 2.8464933 -1.7585610
[51,] 2.7849770 -3.4617375
[52,] 4.3740536 -1.7126827
[53,] 3.8203228 -1.1608206
[54,] 3.5407611 -3.0210995
[55,] 4.3877801 -4.9207311
[56,] 3.3579368 -2.1666392
[57,] 4.2328734 -3.5185138
[58,] 2.5003111 -2.2306257
[59,] 4.9903837 -1.9822267
[60,] 2.2379458 -0.8954690
```

```
plot(x)
```



```
k<- kmeans(x,centers = 2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.221390	-2.563798
2	-2.563798	3.221390

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 81.55993 81.55993
(between_SS / total_SS = 86.0 %)
```

Available components:

[1]	"cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6]	"betweenss"	"size"	"iter"	"ifault"	

Q. from your result object 'k' how many points are in each cluster?

```
k$size
```

```
[1] 30 30
```

Q. What “component” of your results object details the cluster membership?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


q. cluster centers?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. Plot of our cluster

```
plot(x,col=k$cluster)
points(k$centers, col="blue", pch=15, cex=2)
```



We can cluster into 4 groups

```
#kmeans
k4<-kmeans(x, center=4)
#Plot results
plot(x, col=k4$cluster)
```



A big limitation of kmeans is that it does what you ask even if you ask for silly clusters.

Hierarchical Clustering

The main base R function for hierarchical clustering is 'hclust()'. Unlike 'kmeans()' you cant just pass it your data as input. you first need to calculate a distance matrix.

```
d<- dist(x)
hc<-hclust(d)
hc
```

Call:

```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```


use 'plot()' to view results



T

1

N



#Principal Component Analysis (pca)

Here we will do principal component analysis (pca) on some data from the UK.

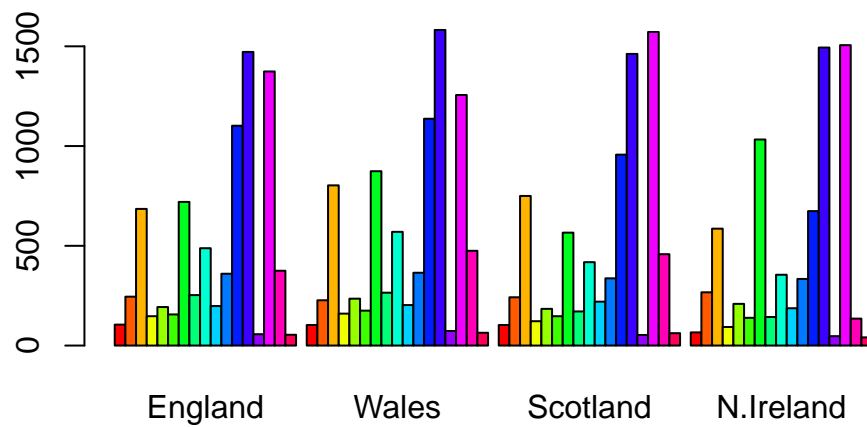
```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

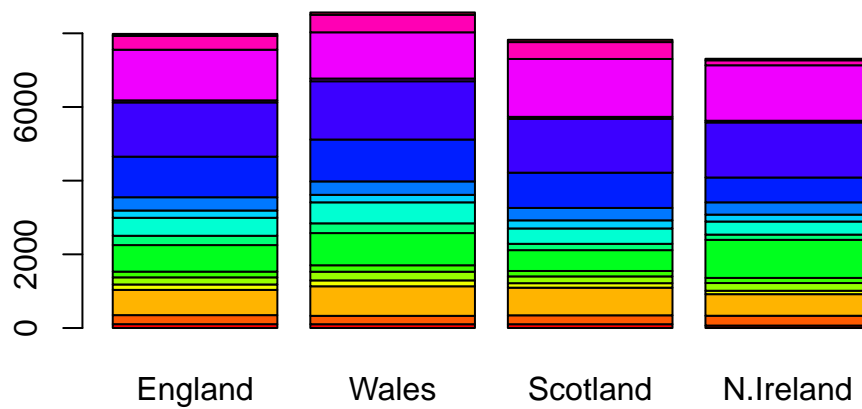
```
#rownames(x)<-x[,-1]
#x<-x[,-1]
#x
```

Spotting differences

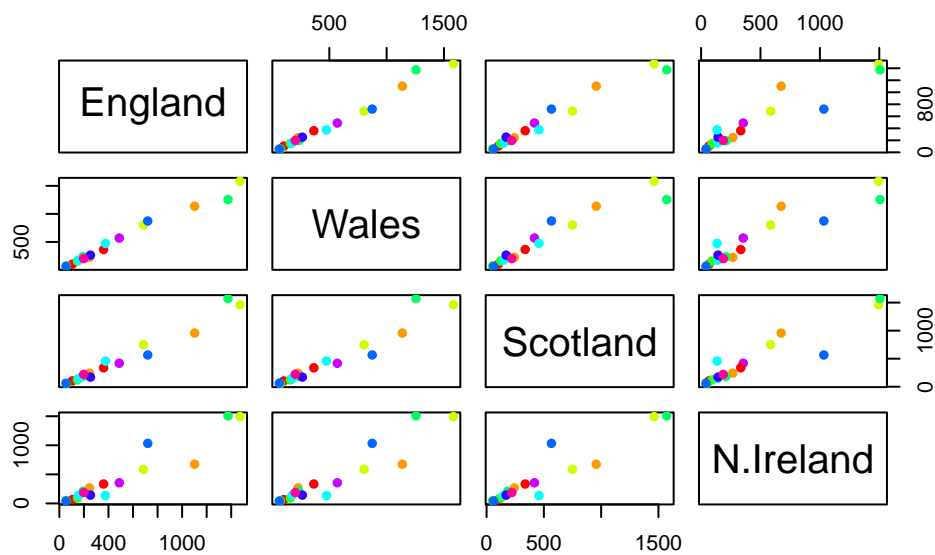
```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



```
pairs(x, col=rainbow(10), pch=16)
```



PCA to the rescue

the main base “r” function for PCA is called ‘prcomp()’. here we need to switch the columns and rows by using ‘t()’

```
pca <-prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Q. How much variance was captured in 2 PCs

96.5%

To make our main “PC score plot” (a.k.a “PC1 VS PC2 plot” or “PC plot” or “ordination plot”)

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class


```
[1] "prcomp"
```

We are after the ‘pca\$x’ result component to make our main PCA plot.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
mycols<-c("orange","red","blue","darkgreen")
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16, xlab = "PC1 (67.4%)", ylab = "PC2 (29%)")
```



Another important result from `pca` is how the original variables (in this case the foods) contribute to the PCs.

This is contained in the `'pca$rotation'` object- often called the “loadings” or “contributions” to the PCs.

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.409382587
Carcass_meat	0.047927628	0.013915823	0.06367111	0.729481922
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.331001134
Fish	-0.084414983	-0.050754947	0.03906481	0.022375878
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.034512161
Sugars	-0.037620983	-0.043021699	-0.03605745	0.024943337
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	0.021396007
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	0.001606882
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.031153231

Processed_potatoes	-0.026886233	0.042850761	-0.07364902	-0.017379680
Processed_Veg	-0.036488269	-0.045451802	0.05289191	0.021250980
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.227657348
Cereals	-0.047702858	-0.212599678	-0.35884921	0.100043319
Beverages	-0.026187756	-0.030560542	-0.04135860	-0.018382072
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.222319484
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.273126013
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001890737

We can make a plot along pc1

```
library(ggplot2)

conrtib<-as.data.frame(pca$rotation)

ggplot(conrtib)+
  aes(PC1, rownames(conrtib))+
  geom_col()
```

